

Oskar Jung

SENIOR DATA ENGINEER

Warthestr. 64, 12051 Berlin, Germany

☎ (+49) 01773711840 | ✉ jung.oskar@gmail.com | 🏠 ojung.io | 📷 ojung | 🌐 oskarjung

Summary

I'm a software engineer with a passion for distributed systems. I enjoy designing, building and operating software solving complex problems. In my career I have helped build data intensive and real-time sensitive applications, as well as platform services.

Skills

Data, Streaming and Backend	Kafka, Hadoop, MapReduce, Hbase, Flink, Vertx, Spark, Redis, Aerospike, Postgres, Node.js
Dev Ops	Kubernetes, Helm, AWS, Terraform, Terragrunt, Packer, Docker, Bash, Unix/Linux, Git
Programming Languages	Java, Scala, Python, Javascript, Clojure

Work Experience

FlixBus Tech GmbH

SENIOR DATA ENGINEER

Berlin, Germany

July 2019 – Present

FlixBus is a German brand that offers intercity bus service in Europe and the United States. It is owned by FlixBus GmbH, which also operates FlixBus and FlixBus. FlixBus Tech is the technology team behind the company.

- Owning Kafka cluster, used as central messaging and streaming platform for 250+ engineers in 50+ technologically heterogeneous teams
- Re-designed existing Kafka cluster to facilitate 45,000 partitions on EC2 with terraform, datadog, lenses, Pagerduty
- Enabling GDPR compliance by integration of Kafka with Azure AD / OpenID Connect
- Design, implementation and deployment of Kafka extensions for authentication, authorization and configuration validation
- Established company-wide Avro schema standard
- Consulting other teams on Kafka and streaming technologies

Ströer Labs Core (ex. mbr targeting GmbH)

SENIOR DATA ENGINEER

Berlin, Germany

Oct 2017 – Jun 2019

Ströer Labs Core (ex. mbr targeting GmbH) is a real-time bidding company. It provides a demand side- and data management platform allowing advertisers to run campaigns on multiple ad exchanges. The company specializes in user-centric optimization using machine learning algorithms. The system scales to above 100k requests per second with a latency of below 100ms.

- Bootstrapped, operated, 250 core, 500Tb Hbase cluster
- Programmatically extended Hbase cluster with coprocessors and filters
- Performance tuned HBase cluster
- Operated 1000 core, petabyte-sized Hadoop cluster on bare metal
- Operated 80 core, 160Tb Kafka cluster with 10Tb of throughput per day
- Implemented streaming Data Management Platform with prediction of user segments based on Kafka, Flink and HBase
- Consulted data scientists and helped them with improving the performance of their ML workflows, as well as with putting them into production

mbr Targeting GmbH

SOFTWARE ENGINEER

Berlin, Germany

Jan 2016 – Sept 2017

- Built a distributed real time bidder using ML approaches to calculate bids for more than 100k events per second with a hard response time requirement of max 100ms
- Performance tuning of Java-based bidder utilizing JMC, JFR and JConsole
- In the process reducing resource consumption by 75% compared to previous implementation

adscale GmbH

SOFTWARE ENGINEER

Christchurch, New Zealand

Apr 2015 – Oct 2015

Adscale Labs is the development team of Adscale GmbH. Adscale is an ad exchange providing a supply side platform to publishers.

- Worked on supply-side real-time-bidding infrastructure handling 140k bids per second
- Worked on video ad server
- Worked on reporting workflows

Education

HTW - Hochschule für Technik und Wirtschaft

B.S. IN APPLIED COMPUTING, GRADE: 2.0 (GOOD)

Berlin, Germany

April 2012 - September 2017

- Thesis: "Implementing a Data Management Platform: Segmentation of Internet Users based on Tracking Data in a Wide Column Database"